

Knowledge Databases Labs

5th Year



matplotlib



jupyter


pandas

Lab-6- Flower Type Prediction with Classification Algorithms

Introduced By:

Eng. Khaled Kondakji Eng. Heba Tadmouri
Eng. Manar AL-Marai

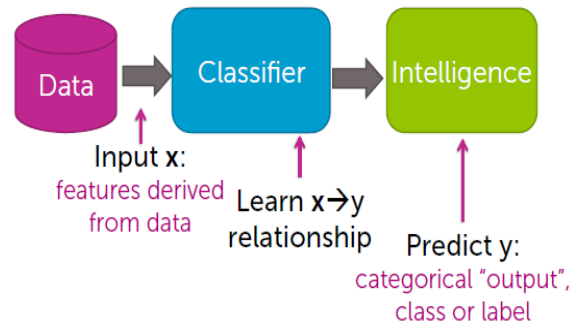
Supervised By:

Dr. Eng. Asmaa Shaar

Flower Type Prediction with Classification Algorithms

Introduction to classification:

- ✎ Given a collection of data tuples (training set). Each tuple is characterized by (x, y):
 - ✓ x (input) is the *attribute set*.
 - ✓ y (output) is the *class label attribute*.
- ✎ Task: based on the input dataset, learn a *classification model (Classifier)* that describes a predetermined set of data classes.



- ✎ Some common algorithms we're going to look at (decision tree, Naïve Bayes and SVM).

Decision Tree Classifier:

- ✎ A decision tree is a flowchart-like tree structure where:
 - ✓ Each internal node denotes a *test on an attribute*.
 - ✓ Each branch represents an *outcome of the test*.
 - ✓ And each leaf node holds a *class label*, the topmost node in the tree is the *root node*.
- ✎ Given a tuple, X, for which the class label is *unknown*,
 - ✓ The *attribute values* of the tuple X are *tested* against the decision tree.
 - ✓ Therefore, a path is traced from the root to a leaf node which holds the class prediction for that tuple.
- ✎ To build up DT we apply the following steps:
 - ✓ A greedy algorithm, where the tree is constructed in a *top-down recursive approach*.
 - ✓ At the start, *all* the training tuples are at the *root node*. Then training set is *recursively* partitioned into smaller subsets as the tree is being built.
 - ✓ training tuples are *partitioned* based on *test attributes* (which are selected based on a *heuristic or statistical measure*), like: (*Information Gain* or *Gain Ratio*).

Naïve Bayes Classifier:

- ✎ Bayesian Classifiers are statistical classifier, performs probabilistic prediction and are based on Bayes' theorem. They have also exhibited high accuracy and speed when applied to large databases.
- ✎ For *classification problems*, we want to determine $P(H|X)$. In other words, we are looking for the probability that tuple x belongs to class c, given that we know the attribute description of x.

Support Vector Machine (SVM) Classifier:


- A linear SVM searches for a hyperplane that maximizes the margin (maximal margin classifier).
- In sklearn we use SVC class which takes kernel as attribute which specifies the SVM kernel like (linear, polynomial, ...).

Flowers' type prediction:

We're going to use Iris dataset that contains information about petal and sepal length and width of different flowers. Our dataset contains the following attributes:

| # | Attribute | Explanation |
|---|---------------------------------|---|
| 1 | Sepal_length | sepal length in cm. |
| 2 | Sepal_width | sepal width in cm. |
| 3 | Petal_length | petal length in cm. |
| 4 | Petal_width | petal width in cm. |
| 5 | species (class label attribute) | Type of flower (Setosa, Versicolor, Virginica). |

Each flower can be one of the following:



Iris Versicolor



Iris Setosa



Iris Virginica

1st step: Import required libraries and load our dataset:

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier
from sklearn import tree
from sklearn.naive_bayes import GaussianNB
from sklearn.svm import SVC
import graphviz
from mlxtend.plotting import plot_decision_regions
from sklearn.metrics import accuracy_score
#load dataset:
iris_df = pd.read_csv('IrisDataset.csv')
iris_df.head()
```

Our dataset looks like the following:

| | sepal_length | sepal_width | petal_length | petal_width | species |
|---|--------------|-------------|--------------|-------------|---------|
| 0 | 5.1 | 3.5 | 1.4 | 0.2 | setosa |
| 1 | 4.9 | 3.0 | 1.4 | 0.2 | setosa |
| 2 | 4.7 | 3.2 | 1.3 | 0.2 | setosa |
| 3 | 4.6 | 3.1 | 1.5 | 0.2 | setosa |
| 4 | 5.0 | 3.6 | 1.4 | 0.2 | setosa |

2nd step: Exploratory Data Analysis (EDA):

First at all we're going to explore our species attribute using `value_counts()` method:

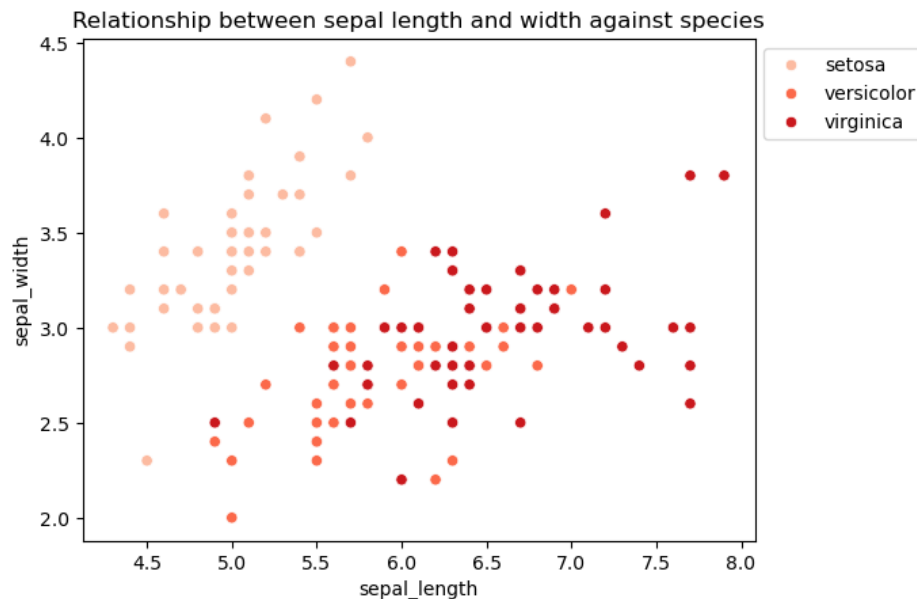
```
iris_df['species'].value_counts()

species
setosa      50
versicolor  50
virginica   50
Name: count, dtype: int64
```

We can extract from here that our classification problem is balanced, meaning that we have quite a similar number of examples in both classes.

Now let's explore relationship between sepal properties and species attribute. For that we're going to use seaborn library:

```
sns.scatterplot(x='sepal_length', y='sepal_width', hue='species',
data=iris_df, palette='Reds')
# Placing Legend outside the Figure
plt.legend(bbox_to_anchor=(1, 1), loc=2)
plt.title('Relationship between sepal length and width against
species')
```

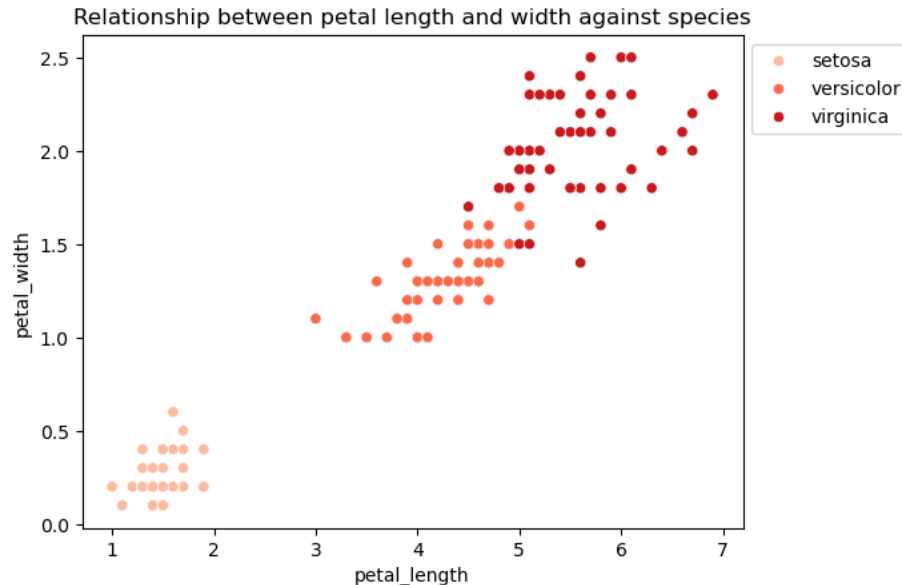


From the above plot, we can infer that:

- 🔍 Species Setosa has smaller sepal lengths but larger sepal widths.
- 🔍 Versicolor Species lies in the middle of the other two species in terms of sepal length and width.
- 🔍 Species Virginica has larger sepal lengths but smaller sepal widths.

In the same way we're going to draw the relationship between petal properties and species attribute

```
sns.scatterplot(x='petal_length', y='petal_width', hue='species', data=iris_df,
palette='Reds')
# Placing Legend outside the Figure
plt.legend(bbox_to_anchor=(1, 1), loc=2)
plt.title('Relationship between petal length and width against species')
```




From the above plot, we can infer that:

- 🔧 Species Setosa has smaller petal lengths and widths.
- 🔧 Versicolor Species lies in the middle of the other two species in terms of petal length and width.
- 🔧 Species Virginica has the largest of petal lengths and widths.

Another way to get insights about data attributes is by using `histplot()` which is used basically for the univariant set of observations and visualizes it through a histogram i.e. only one observation and hence we choose one particular column of the dataset.

```
g = sns.FacetGrid(iris_df.melt(id_vars='species'), col="variable", sharex=False,
sharey=False, hue="species")
g.map(sns.histplot, "value")
```



Refer to this [link](#) for information. From the above plots, we can see that:

- 🔧 In the case of Sepal Length, there is a huge amount of overlapping.
- 🔧 In the case of Sepal Width also, there is a huge amount of overlapping.
- 🔧 In the case of Petal Length, there is a very little amount of overlapping.
- 🔧 In the case of Petal Width also, there is a very little amount of overlapping.

Finally, we're going to print the correlation matrix between attributes:

```
sns.heatmap(iris_df.corr(numeric_only=True), annot = True, cmap='Reds');
plt.show()
```



3rd step: Prepare our dataset for classification task:

```
# shuffle our data:
iris_df = iris_df.sample(frac=1) #frac = 1 mean returns all rows.

# extract input and output label:
X = iris_df.drop(labels=["species"],axis = 1)
y = iris_df['species']

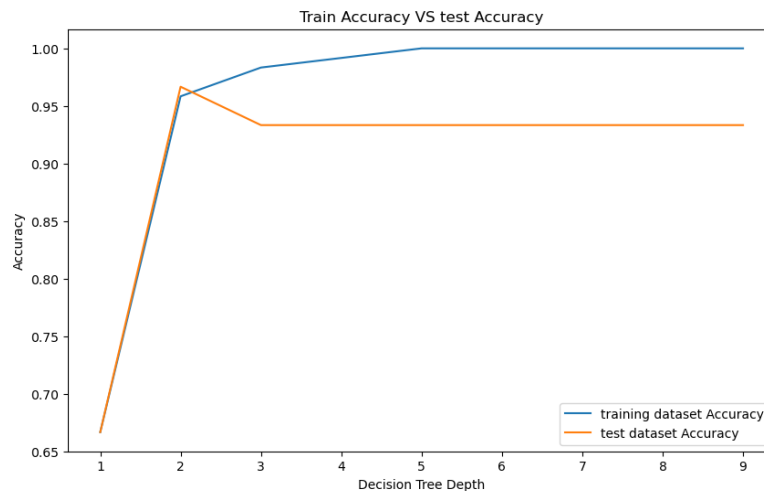
#setting a seed and split our data into train and test sets:
np.random.seed(42)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2)
```

4th step: build and tune our classification models:

Build and evaluate Decision Tree Model:


We're going to use `DecisionTreeClassifier` class provided by `sklearn` in order to learn a model based on the input data. Our class takes some important parameters such as `criterion` and `max_depth`.

```
train_acc = [] #training set accuracy list
test_acc = [] #test set accuracy list
train_error = [] #training set error list
test_error= [] #test set error list
depths = list(range(1,10))
for i in range (1,10):
    dt = DecisionTreeClassifier(criterion='entropy',max_depth=i)
    dt.fit(X_train,y_train)
    train_acc.append(accuracy_score(y_train,dt.predict(X_train)))
    test_acc.append(accuracy_score(y_test,dt.predict(X_test)))
    train_error.append(1-(accuracy_score(y_train,dt.predict(X_train))))
    test_error.append(1-(accuracy_score(y_test,dt.predict(X_test))))
plt.figure(figsize=(10,6))
plt.plot(depths, train_acc,'-',depths,test_acc,'-')
plt.title("Train Accuracy VS test Accuracy")
plt.xlabel("Decision Tree Depth")
plt.ylabel("Accuracy")
plt.legend(["training dataset Accuracy", "test dataset Accuracy"])
```



- 🔧 We can notice that if we use the *training set* to measure the classifier's accuracy, this estimate would likely be *optimistic* because the classifier tends to *overfit the data*.
- 🔧 Because of that, we tend to estimate classifier accuracy using a list of tuples called *test dataset* which is *independent* of the training dataset.
- 🔧 The *accuracy* of a classifier, M, on a given test set is the "percentage of test set tuples that are *correctly* classified by the classifier".
- 🔧 We can evaluate the same concepts based on classification error which is given by:
 - Error = 1- Accuracy.

```
plt.figure(figsize=(10,6))
plt.plot(depths, train_error, '-', depths, test_error, '-')
plt.title("Train Error VS test Error")
plt.xlabel("Decision Tree Depth")
plt.ylabel("Error Rate")
plt.legend(["training dataset Error", "test dataset Error"])
```

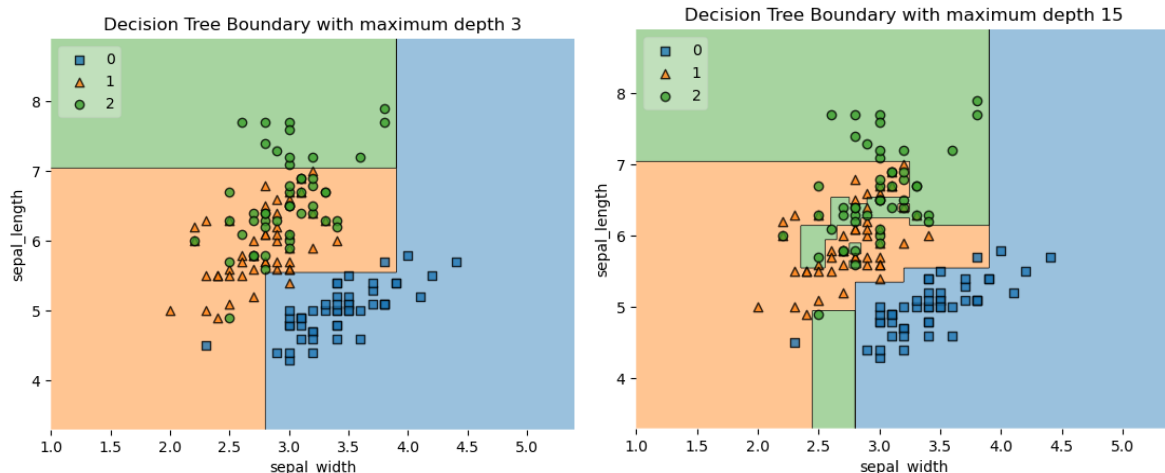


- 🔧 To select the best `max_depth` value for a decision tree, we should notice the effect of getting higher depth trees VS lower depth trees.
- 🔧 Now, we're going to build a new `DecisionTreeClassifier` based only on two attributes (`sepal width`, `sepal length`), in order to notice the `max_depth` parameter effect for 2 values (3 and 15).

```

x1 = iris_df[["sepal_width","sepal_length"]].values
y = iris_df.species.replace({'setosa':0,'versicolor':1,'virginica':2}).values
x1_train, x1_test, y1_train, y1_test = train_test_split(x1,y,test_size= .2)
for i in [3,15]:
    dt1 = DecisionTreeClassifier(criterion='entropy',max_depth=i)
    dt1.fit(x1_train,y1_train)
    plot_decision_regions(np.array(x1, ), np.array(y), clf=dt1, legend=2,)
    plt.xlabel('sepal_width')
    plt.ylabel('sepal_length')
    title = 'Decision Tree Boundary with maximum depth '+ str(i)
    plt.title(title)
    plt.show()

```



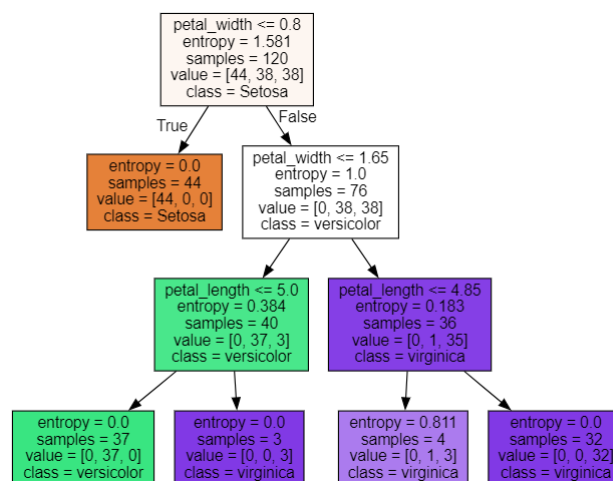
When a decision tree is built, many of the branches will reflect anomalies in the training data due to *noise* or *outliers*. Pruned trees:

- ✓ Tend to be smaller and *less complex*, and thus *easier* to comprehend. Usually faster and *better* at *classifying test tuples* than unpruned trees. This code will draw the fitted tree:

```

dt = DecisionTreeClassifier(criterion='entropy', max_depth=3)
dt.fit(X_train,y_train)
dot_data=tree.export_graphviz(dt,out_file=None,feature_names=list(X.columns),class_names=['Setosa', 'versicolor', 'virginica'], filled = True)
tr = graphviz.Source(dot_data, format = "png")
tr

```



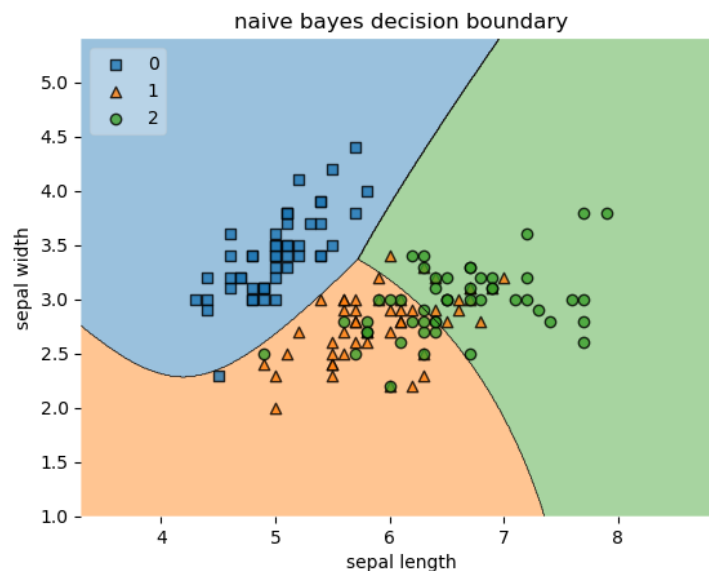
Build model using naïve bayes model:

SKLearn provides many classes that implements naïve bayes classifier like GaussianNB, MultinomialNB, ... which can be used for different purposes.

```
classifier = GaussianNB() #create object from GaussianNB class
classifier.fit(X_train, y_train) #fit model to dataset
y_pred = classifier.predict(X_test) #predict test set using fitted model
print ("Accuracy: ", accuracy_score(y_test, y_pred)) #measure model accuracy
```


again, we're going to use only sepal length and width for fit a gaussian naïve bayes model. We also replace our classes values with numeric values to draw the boundary like follows:

```
X_data = iris_df.iloc[:,0:2].values
y_labels =
iris_df.species.replace({'setosa':0,'versicolor':1,'virginica':2}).copy().values
nb = GaussianNB().fit(X_data,y_labels)
fig = plot_decision_regions(X=X_data, y=y_labels,clf=nb, legend=2)
plt.title('naïve bayes decision boundary')
plt.xlabel('sepal length')
plt.ylabel('sepal width')
```



Support Vector Machine (SVM) Model:

```
model = SVC(kernel='linear')
model.fit(X_train,y_train)
print ("Accuracy : ", accuracy_score(y_test, model.predict(X_test)))
#draw boundaries
model1 = SVC(kernel='linear')
model1.fit(X_data ,y_labels)
fig = plot_decision_regions(X=X_data, y=y_labels,clf=model1, legend=2)
plt.title('SVM decision boundary')
plt.xlabel('sepal length')
plt.ylabel('sepal width')
```



References:

- 1- <https://www.geeksforgeeks.org/exploratory-data-analysis-on-iris-dataset/>
- 2- Mlxtend docs: <http://rasbt.github.io/mlxtend/>
- 3- <https://towardsdatascience.com/machine-learning-basics-naive-bayes-classification-964af6f2a965>
- 4- <https://www.pycodemates.com/2022/10/classification-of-iris-dataset-using-SVM-in-python.html>
- 5- Pandas docs <https://pandas.pydata.org/docs/> .
- 6- NumPy docs <https://numpy.org/doc/> .
- 7- Matplotlib docs <https://matplotlib.org/3.3.3/contents.html> .
- 8- SciKit-learn docs <https://scikit-learn.org/0.21/documentation.html> .
- 9- http://rasbt.github.io/mlxtend/user_guide/plotting/plot_decision_regions/#example-1-decision-regions-in-2d